

Stan interface

Week #8

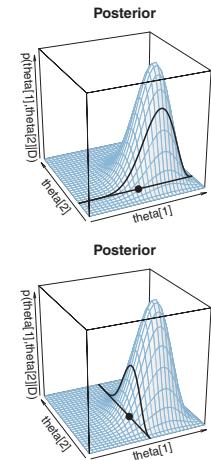
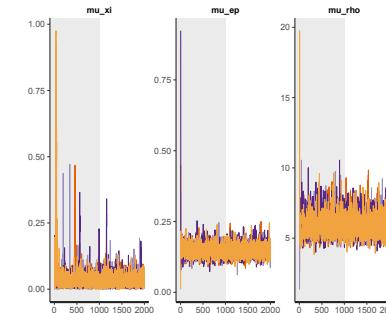


*Woo-Young Ahn
Department of Psychology
Seoul National University
ccs-lab.github.io*

Recap of Week 7

- *Markov Chain Monte Carlo*
 - *What is a Markov Chain?*
 - *What is a Monte Carlo process?*
- *Metropolis algorithm*
 - *How does it work? E.g., travelling politician*
 - *What are pros and cons?*
- *Gibbs sampling*
 - *How does it work?*
 - *What are pros and cons?*
- *Things to know when using MCMC*
 - *Traceplots (“fat hairy caterpillar”)*
 - *Rhat*
 - *Why multiple chains?*
 - *Burn-in (warmup) samples*
 - *Autocorrelation, effective sample size, and thinning*

$$p_{\text{move}} = \min \left(\frac{P(\theta_{\text{proposed}})}{P(\theta_{\text{current}})}, 1 \right)$$



Bayesian Demo

- <https://mcl.shinyapps.io/prior2post/>

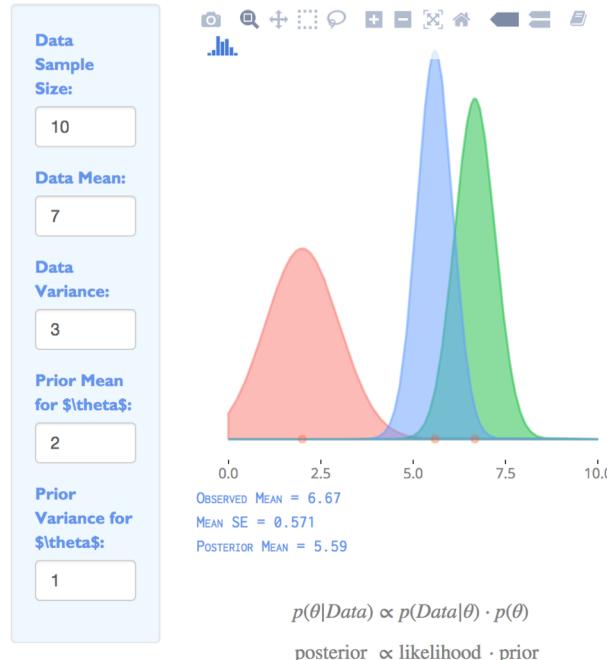
Bayesian Demo

The following demo regards estimating a mean (ignoring estimating the variance) and how the prior distribution and likelihood combine to produce the posterior distribution.

You can set the following parameters:

- Sample size: 0-250
- Observed and Prior Means: 1-10
- Observed and Prior Variances: 1-5

The tested θ parameters are a sequence of 500 values from 0 to 10.



More MCMC algorithms

- <http://m-clark.github.io/docs/lmcmc>
- *Dozens of MCMC algorithms covered*

The screenshot shows a red header bar with the date "2016-11-27" and the title "MCMC Algorithms". Below the header is a navigation menu with links to "Preface", "Markov Chain Monte Carlo", and a list of algorithms. The main content area lists over 40 different MCMC algorithms, each preceded by a small blue square icon.

2016-11-27

MCMC Algorithms

Preface

Markov Chain Monte Carlo

Adaptive Directional Metropolis-within-Gibbs

Automated Factor Slice Sampler

Adaptive Griddy-Gibbs

Adaptive Hamiltonian Monte Carlo

Affine-Invariant Ensemble Sampler

Adaptive Metropolis

Adaptive-Mixture Metropolis

Adaptive Metropolis-within-Gibbs

Componentwise Hit-And-Run Metropolis

Differential Evolution Markov Chain

Delayed Rejection Adaptive Metropolis

Delayed Rejection Metropolis

Elliptical Slice Sampler

Algorithms

The MCMC algorithms in the `Laplace'sDemon` package are now presented alphabetically:

- Adaptive Directional Metropolis-within-Gibbs
- Adaptive Griddy-Gibbs
- Adaptive Hamiltonian Monte Carlo
- Adaptive Metropolis
- Adaptive Metropolis-within-Gibbs
- Adaptive Mixture Metropolis
- Adaptive Metropolis
- Affine-Invariant Ensemble Sampler
- Automated Factor Slice Sampler
- Componentwise Hit-And-Run Metropolis
- Delayed Rejection Adaptive Metropolis
- Delayed Rejection Metropolis
- Differential Evolution Markov Chain
- Elliptical Slice Sampler
- Gibbs Sampler
- Griddy-Gibbs
- Hamiltonian Monte Carlo
- Hamiltonian Monte Carlo with Dual-Averaging
- Hit-And-Run Metropolis
- Independence Metropolis
- Interchain Adaptation
- Metropolis-Adjusted Langevin Algorithm
- Metropolis-Coupled Markov Chain Monte Carlo
- Metropolis-within-Gibbs
- Multiple-Try Metropolis
- No-U-Turn Sampler
- Preconditioned Crank-Nicolson
- Oblique Hyperrectangle Slice Sampler
- Random Dive Metropolis-Hastings
- Random-Walk Metropolis
- Reflective Slice Sampler
- Refractive Sampler
- Reversible-Jump
- Robust Adaptive Metropolis
- Sequential Adaptive Metropolis-within-Gibbs
- Sequential Metropolis-within-Gibbs
- Slice Sampler
- Stochastic Gradient Langevin Dynamics
- Tempered Hamiltonian Monte Carlo
- t-walk
- Univariate Eigenvector Slice Sampler
- Updating Sequential Adaptive Metropolis-within-Gibbs
- Updating Sequential Metropolis-within-Gibbs

Installation (rstan)

- *For MAC users:*
 1. *Make sure you have Xcode installed:*
<https://developer.apple.com/xcode/>
- *For Windows users:*
 1. *Download R tools:* <https://cran.r-project.org/bin/windows/Rtools/>
 - *For more details, see:* <https://github.com/stan-dev/rstan/wiki/Installing-RStan-on-Windows>

Run the following command through the R console:

```
install.packages("rstan", dependencies = T)
```

Stan

Stanislaw Ulam, namesake of Stan and co-inventor of Monte Carlo methods ([Metropolis and Ulam, 1949](#)), shown here holding the Fermiac, Enrico Fermi's physical Monte Carlo simulator for neutron diffusion.

Image from ([Giesler, 2000](#)).



Stan

- *A modeling language (software) for Bayesian data analysis*
- *MCMC is actually done in C++*
- *Stan interfaces in R, Python, Matlab, Julia, Stata, Mathematica, and Scala*
- *Hamiltonian Monte Carlo (HMC), which works well even for complicated models and models with highly correlated parameters*
- *Also supports variational Bayes*
- <http://mc-stan.org/>
- *Maintained by Andrew Gelman and his team at Columbia University*

Stan

- *Manual (version 2.17.0)*
 - <https://github.com/stan-dev/stan/releases/download/v2.17.0/stan-reference-2.17.0.pdf>
 - Over 600 pages long...
- *Stan forum:* <http://discourse.mc-stan.org/>
- *Similar to BUGS (and JAGS)*
 - See Appendix B of Stan manual

Grammar in Stan

- Use `//` and `/* */` for comments
- Every assignment statement must be followed by a semicolon (`;`)
e.g., `n = n + 1;`
- Use `=` instead of `<-` for assignment
- ...

Structure of Stan model specification

- *6 model blocks*

```
data {  
  ... read in external data ...  
}  
transformed data {  
  ... pre-processing of data ...  
}  
parameters {  
  ... parameters to be sampled by HMC ...  
}  
transformed parameters {  
  ... pre-processing of parameters ...  
}  
model {  
  ... statistical/cognitive model ...  
}  
generated quantities {  
  ... post-processing of the model ...  
}
```

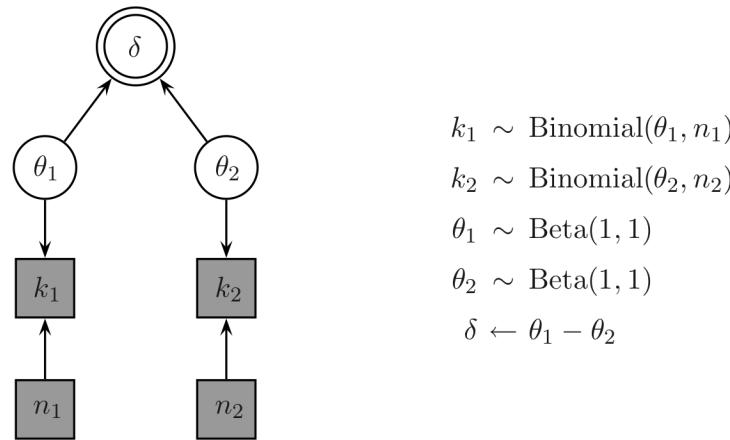
Optional

Group discussion

- *Discuss what should be in each block in Stan.*
- *Explain to each other what's unique and essential part of each block.*
- *Look up Stan manual if needed.*

[https://github.com/stan-dev/stan-releases/download/v2.17.0/stan-reference-2.17.0.pdf](https://github.com/stan-dev/stan/releases/download/v2.17.0/stan-reference-2.17.0.pdf)

Graphical model



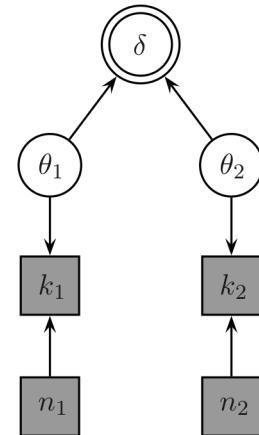
Variables without shading: unobserved
Variables with shading: observed

Continuous variables: circular nodes
Discrete variables: square nodes

Single vs double bordered: probabilistic vs deterministic

1. Data block in Stan

- We tell Stan which data should be loaded
- Data types
 - Real (real)
 - Integer (int)
 - Vector and matrix types
 - `vector` for column vectors
 - `row_vector` for row vectors
 - `matrix` for matrices
- Constrained data types
 - Data types may be given lower and upper bounds
 - e.g., `vector<lower = -1, upper = 1>[3] rho;`



Conditional operators (page 62 of Stan manual)

<i>Op.</i>	<i>Prec.</i>	<i>Assoc.</i>	<i>Placement</i>	<i>Description</i>
? :	10	right	ternary infix	conditional
	9	left	binary infix	logical or
&&	8	left	binary infix	logical and
==	7	left	binary infix	equality
!=	7	left	binary infix	inequality
<	6	left	binary infix	less than
<=	6	left	binary infix	less than or equal
>	6	left	binary infix	greater than
>=	6	left	binary infix	greater than or equal
+	5	left	binary infix	addition
-	5	left	binary infix	subtraction
*	4	left	binary infix	multiplication
/	4	left	binary infix	(right) division
%	4	left	binary infix	modulus
\	3	left	binary infix	left division
.*	2	left	binary infix	elementwise multiplication
./	2	left	binary infix	elementwise division
!	1	n/a	unary prefix	logical negation
-	1	n/a	unary prefix	negation
+	1	n/a	unary prefix	promotion (no-op in Stan)
^	0.5	right	binary infix	exponentiation
,	0	n/a	unary postfix	transposition
O	0	n/a	prefix, wrap	function application
[]	0	left	prefix, wrap	array, matrix indexing

Indexing in Stan (page 65 of Stan manual)

<i>index type</i>	<i>example</i>	<i>value</i>
integer	<code>a[11]</code>	value of <code>a</code> at index 11
integer array	<code>a[ii]</code>	<code>a[ii[1]], ..., a[ii[K]]</code>
lower bound	<code>a[3:]</code>	<code>a[3], ..., a[N]</code>
upper bound	<code>a[:5]</code>	<code>a[1], ..., a[5]</code>
range	<code>a[2:7]</code>	<code>a[2], ..., a[7]</code>
all	<code>a[:]</code>	<code>a[1], ..., a[N]</code>
all	<code>a[]</code>	<code>a[1], ..., a[N]</code>

Figure 4.2: *Types of indexes and examples with one-dimensional containers of size N and an integer array ii of type int[] size K.*

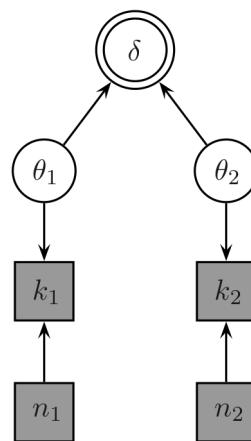
2. Transformed Data block in Stan

- *Sometimes we need to transform data loaded in the Data block before we feed the (transformed) data into our model*
- *This is optional. Often easier to do it in R first.*
- *E.g.,*

```
transformed data {          // Transformed data block
    vector[N] logX;
    logX = log(X);
}
```

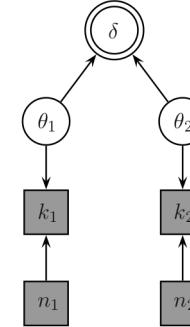
3. Parameters & transformed parameters blocks in Stan

- We define parameters to be sampled
- Transformed parameters → parameters that are dependent on the parameters



4. Model block

- *Specify priors and likelihood*
- *Declare any variables necessary*
- *Look up Stan manual and find functions/commands for each prior and transformation*
- *You can use for loop, if statement, etc.*
 - *(not just for model but for other blocks as well)*



$$\begin{aligned}k_1 &\sim \text{Binomial}(\theta_1, n_1) \\k_2 &\sim \text{Binomial}(\theta_2, n_2) \\\theta_1 &\sim \text{Beta}(1, 1) \\\theta_2 &\sim \text{Beta}(1, 1) \\\delta &\leftarrow \theta_1 - \theta_2\end{aligned}$$

5. Generated quantities block

- *We can generate some sorts of quantities based on already sampled variables*
- *Posterior predictive checks*
- *Log probability for a model*

Distributions and functions used in Stan

- *We need to look up Stan manual*
- *Distributions → part VIII (pages 502~)*

54.1. Normal Distribution

Probability Density Function

If $\mu \in \mathbb{R}$ and $\sigma \in \mathbb{R}^+$, then for $y \in \mathbb{R}$,

$$\text{Normal}(y|\mu, \sigma) = \frac{1}{\sqrt{2\pi} \sigma} \exp\left(-\frac{1}{2} \left(\frac{y-\mu}{\sigma}\right)^2\right).$$

Sampling Statement

```
y ~ normal(mu, sigma);
```

Increment log probability with `normal_lpdf(y | mu, sigma)`, dropping constant additive terms; Section 5.3 explains sampling statements.

Distributions and functions used in Stan

- *Built-in functions (pages 427-)*

R **inv_logit**(T x)

Returns the (elementwise) logistic sigmoid function applied to x,

$$\text{inv_logit}(x) = \frac{1}{1 + \exp(-x)}$$

for any argument type T; see Section 41.1 for details including return type R.

Example

ra_prospect_stan_singleSubj

- *Let's look at the code together!*